

Daniel Nakashima => cs61c-hi

Chung Fai Tsim => cs61c-eh

Matrix Multiply - Project 3 - Writeup Part 2

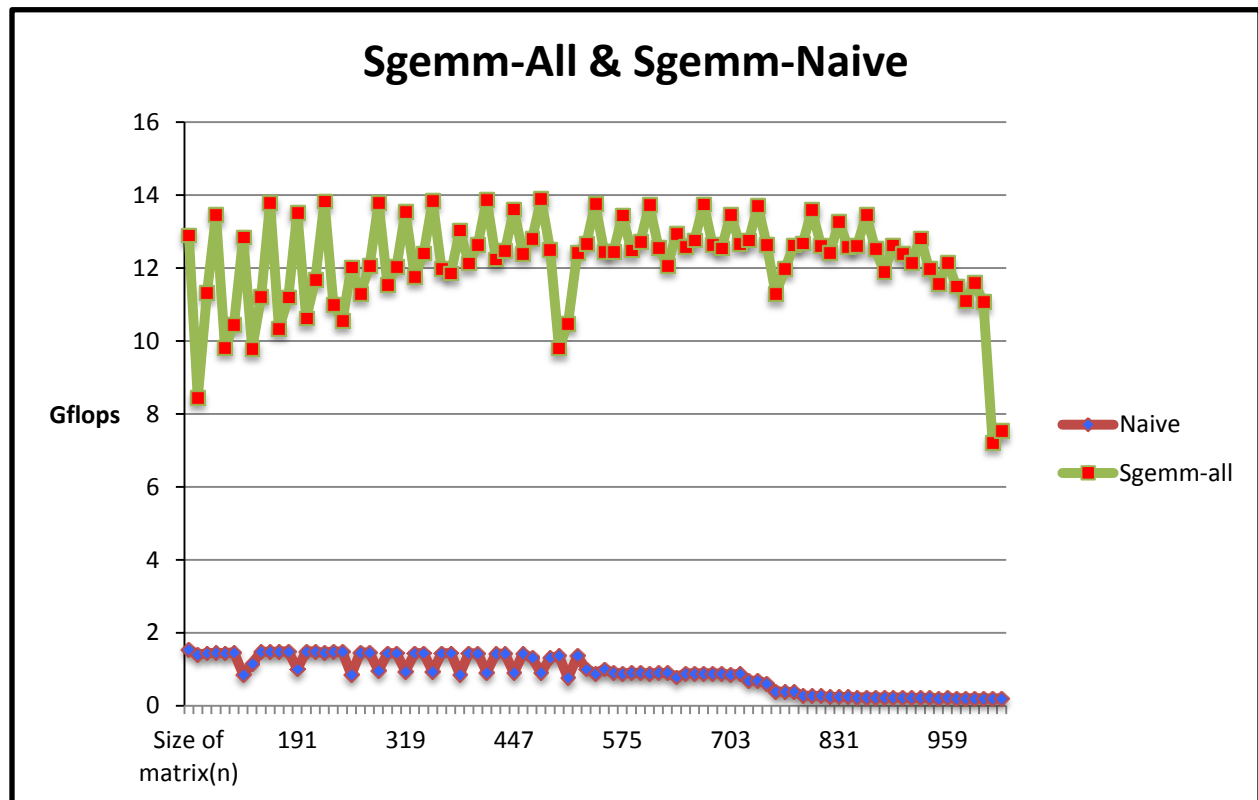
1. A brief description of any changes you made to your code from part 1 in order to get it to run well on a range of matrix sizes (max 150 words)

In our first submission, in the event that input size of the matrix was not a multiple of 16, we handled this naively with a single intrinsic and purely naively for the case for the fringes. In our sgemm-all submission, our approach was to “pad” our input matrix size N that was not divisible by 16 to the ideal size. IE, $N = N/8*8 + 8$. To do this, we iteratively pad the each of the A , B and C matrices with zeros and from there perform the same operations that we did for sgemm-small i.e. ideal sizes. This especially improved the performance for the matrix sizes that are not divisible by 8.

2. A brief description of how you used OpenMP pragmas to parallelize your code in sgemm-openmp.c (max 150 words).

In sgemm-all.c, the nested loop are in the order of k, j, i . We stored the result of $C = C + A*B$ back in the most inner loop i according to $C = k+j*n$. Since we switched to use openmp, we need to split each of the work equivalently into 8 cores. So what we need to do is switch the loop order of k and j . In our sgemm-openmp.c, we have the loop order j, k, i . For example, if $N=64$ and we have 8 cores, then cpu0 is calculating the result and store in the column 0-7 in Matrix C , and cpu7 is calculating the result and store in the column 56-63 in Matrix C . We set every variable we used inside the nested loop to be private except matrix size N .

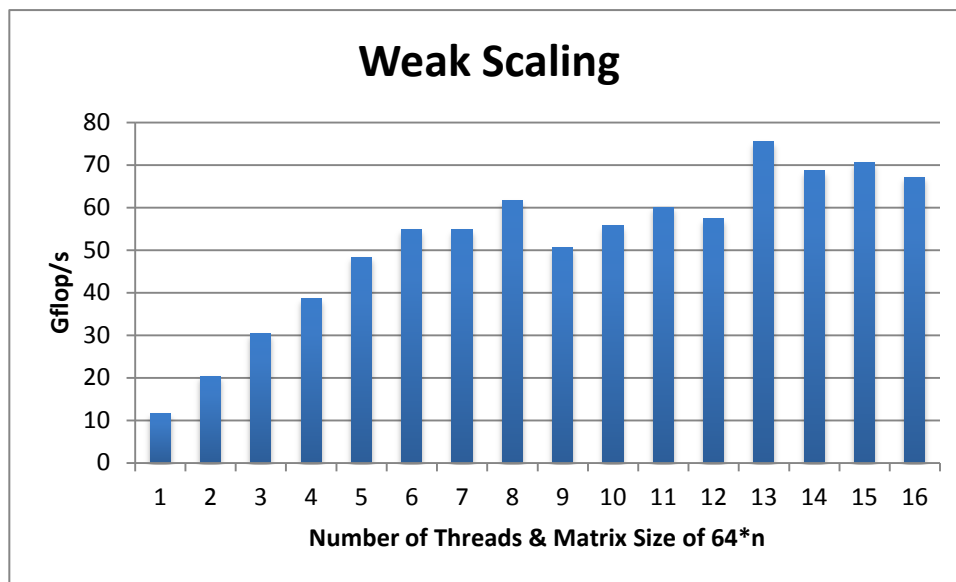
3. A plot showing the speedup of sgemm-all.c over sgemm-naive.c for values of N between 64 and 1024



4. A weak scaling plot of the performance of your sgemm-openmp.c code (use your sgemm-all.c code as the baseline for the single threaded case)

Number of threads range = 1-16

Matrix Size => multiples of 64 i.e. 2 threads ran a matrix size of $2 \times 64 = 128$



5. A strong scaling plot of the performance of your sgemm-openmp.c code

Used fixed matrix size $n = 256$;

Number of threads range = 1 - 16;

Strong Scaling

